

INSTITUTO TECNOLÓGICO DE AERONÁUTICA



Luis Guilherme Gomes Aguiar

**SIMULATED HUMANOID ROBOT CONTROL
WITH REINFORCEMENT LEARNING**

Final Paper
2018

Course of Electronics Engineering

Luis Guilherme Gomes Aguiar

**SIMULATED HUMANOID ROBOT CONTROL
WITH REINFORCEMENT LEARNING**

Advisor

Prof. Dr. Takashi Yoneyama (ITA)

Co-advisor

Prof. Dr. Marcos Ricardo Omena de A. Máximo (ITA)

ELECTRONICS ENGINEERING

SÃO JOSÉ DOS CAMPOS
INSTITUTO TECNOLÓGICO DE AERONÁUTICA

2018

Cataloging-in Publication Data
Documentation and Information Division

Aguiar, Luis Guilherme Gomes
Simulated Humanoid Robot Control with Reinforcement Learning / Luis Guilherme Gomes Aguiar.
São José dos Campos, 2018.
42f.

Final paper (Undergraduation study) – Course of Electronics Engineering– Instituto Tecnológico de Aeronáutica, 2018. Advisor: Prof. Dr. Takashi Yoneyama. Co-advisor: Prof. Dr. Marcos Ricardo Omena de A. Máximo.

1. Dinamica de robos. 2. Robos humanoides. 3. Controle de robos. 4. Inteligencia artificial. 5. Robotica. 6. Controle. I. Instituto Tecnológico de Aeronáutica. II. Title.

BIBLIOGRAPHIC REFERENCE

AGUIAR, Luis Guilherme Gomes. **Simulated Humanoid Robot Control with Reinforcement Learning**. 2018. 42f. Final paper (Undergraduation study) – Instituto Tecnológico de Aeronáutica, São José dos Campos.

CESSION OF RIGHTS

AUTHOR'S NAME: Luis Guilherme Gomes Aguiar

PUBLICATION TITLE: Simulated Humanoid Robot Control with Reinforcement Learning.

PUBLICATION KIND/YEAR: Final paper (Undergraduation study) / 2018

It is granted to Instituto Tecnológico de Aeronáutica permission to reproduce copies of this final paper and to only loan or to sell copies for academic and scientific purposes. The author reserves other publication rights and no part of this final paper can be reproduced without the authorization of the author.

Luis Guilherme Gomes Aguiar
Rua H8A, 131
12228-460 – São José dos Campos-SP

SIMULATED HUMANOID ROBOT CONTROL WITH REINFORCEMENT LEARNING

This publication was accepted like Final Work of Undergraduation Study

Luis Guilherme Gomes Aguiar
Author

Takashi Yoneyama (ITA)
Advisor

Marcos Ricardo Omena de A. Máximo (ITA)
Co-advisor

Prof. Dr. Cairo Nascimento
Course Coordinator of Electronics Engineering

São José dos Campos: Junho 26, 2018.

To God, my parents and all my good
friends.

Acknowledgments

make it later

Resumo

Abstract

List of Figures

FIGURE 1.1 – DeepMind recent achievements.	16
FIGURE 1.2 – Simulated humanoid agent movements and AI.	17
FIGURE 1.3 – Robocup symbol and Soccer 3D Simulation league match.	17
FIGURE 2.1 – Agent interacting with environment.	20
FIGURE 2.2 – Example of board games that have large discrete state spaces. . . .	24
FIGURE 2.3 – Sample control tasks from the (BROCKMAN <i>et al.</i> , 2016) toolkit. . . .	24
FIGURE 2.4 – Actor-critic architecture from (SUTTON; BARTO, 1998).	26
FIGURE 3.1 – Neuron unit cell structure.	29
FIGURE 3.2 – Simple neural network.	30
FIGURE 3.3 – Node computation.	31
FIGURE 3.4 – Linear rectified unit plot.	31
FIGURE 3.5 – Neural network with multiple layer.	34
FIGURE 3.6 – Gradient descent on function g	36
FIGURE A.1 – Uma figura que está no apêndice	41

List of Tables

List of Abbreviations and Acronyms

AI	Artificial Intelligence
API	Application Programming Interface
ANN	Artificial Neural Network
CMA-ES	Covariance Matrix Adaptation Evolution Strategy
CL	Curriculum Learning
CoM	Center of Mass
CNN	Convolutional Neural Network
DNN	Deep Neural Network
DDPG	Deep Deterministic Policy Gradients
DPG	Deterministic Policy Gradients
DQN	Deep Q-Network
DRL	Deep Reinforcement Learning
GPU	Graphics Processing Unit
KL	Kullback-Leibler
IK	Inverse Kinematics
MC	Monte-Carlo
MDP	Markov Decision Process
ML	Machine Learning
MLP	Multilayer Perceptron
NN	Neural Network
PPO	Proximal Policy Optimization
ReLU	Rectified Linear Unit
RL	Reinforcement Learning
RNN	Recurrent Neural Network
RPC	Remote Procedure Call
SS3D	Soccer Simulation 3D
Soccer3D	Soccer Simulation 3D
SGD	Stochastic Gradient Descent
TD	Temporal-Difference

List of Symbols

\mathbb{A}	A set
\mathbb{R}	The set of real numbers
a_i	Element i of vector a , with indexing starting at 1
A^T	Transpose of matrix A
$\frac{dy}{dx}$	Derivative of y with respect to x
$\frac{\partial y}{\partial x}$	Partial derivative of y with respect to x
∇y	Gradient of y
$\int_a^b f(x)dx$	Definite integral with respect to x over $[a, b]$
$\mathbb{I}^{\text{condition}}$	conditional unit function
$f : \mathbb{A} \rightarrow \mathbb{B}$	The function f with domain \mathbb{A} and range \mathbb{B}
$f(x, \theta)$	A function of x parameterized by θ
$\mathbb{P}(x)$	Probability distribution over a continuous variable
$\mathbb{E}[X]$	Expectation of random variable X
$\text{KL}[P, Q]$	KL divergence between P and Q
$\log x$	Natural logarithm of x
Reinforcement Learning	
s, s'	States
a	Action
r	Reward
t	Discrete timestep
A_t	Action at timestep t
S_t	State at timestep t
R_t	Reward at timestep t
G_t	Return (cumulative discounted reward) following time t
π	Policy, agent behavior rule
$\pi(s)$	Action taken in state s under policy π

$v_\pi(s)$	Value of state s under policy π (expected return)
$v_*(s)$	Value of state s under the optimal policy
$q_\pi(s, a)$	Value of taking action a in state s under policy π
$q_\pi(s, a)$	Value of taking action a state s under policy π
V, V_t	Estimate of state-value function v_π or v_*
Q, Q_t	Estimate of action-value function q_π or q_*

Contents

1	INTRODUCTION	15
1.1	Motivation	15
1.2	Problem Statement	17
1.3	Approach	18
1.4	Literature Review	18
1.5	Contributions	19
1.6	Outline of this Dissertation	19
2	REINFORCEMENT LEARNING	20
2.1	Markov Decision Process	21
2.2	Return	21
2.3	Policies	22
2.4	Value functions	22
2.5	Goal of RL	22
2.5.1	Optimal Value Function	22
2.5.2	Optimal Policy	23
2.6	Function Approximation	23
2.7	RL algorithms	25
2.7.1	Value Function Methods	25
2.7.2	Policy Search Methods	25
2.7.3	Actor-Critic Methods	26
2.8	Curriculum Learning	27
3	DEEP LEARNING	28

3.1	History	28
3.2	Neural Networks	29
3.2.1	Representation	29
3.2.2	Vectorization	32
3.2.3	Deep Neural Networks	33
3.3	Learning	33
3.3.1	Cost Function	34
3.3.2	Optimization Algorithms	35
3.4	Back Propagation	37
BIBLIOGRAPHY		38
APPENDIX A – TÏ>½PICOS DE DILEMA LINEAR		41
A.1	Uma Primeira Seção para o Apêndice	41
ANNEX A – EXEMPLO DE UM PRIMEIRO ANEXO		42
A.1	Uma Seção do Primeiro Anexo	42

1 Introduction

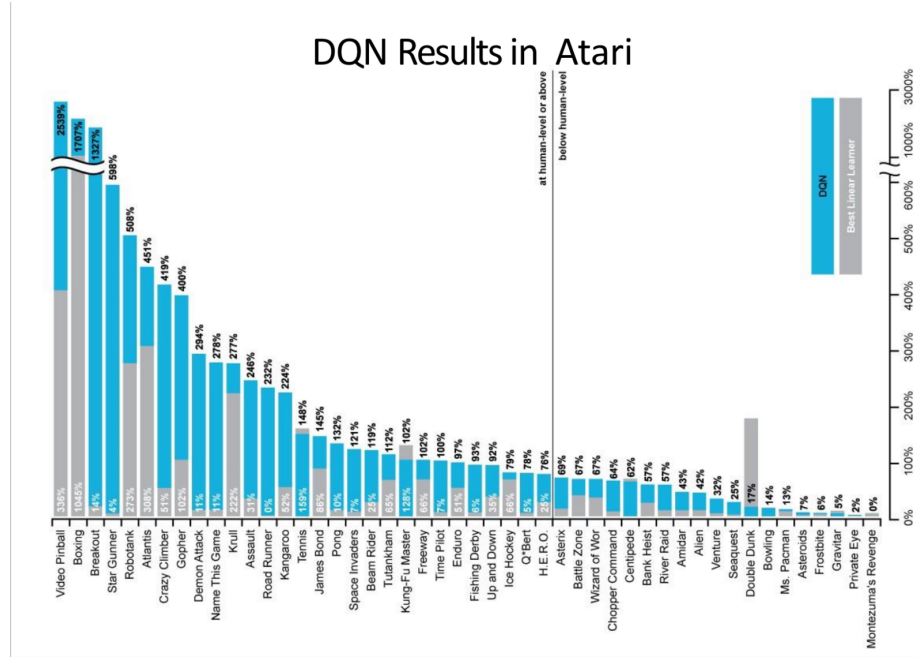
1.1 Motivation

In the last decades, the advances in computers architecture and computer science have pushed the research frontier to produce super intelligent algorithms, capable of dealing with difficult classification problems of subtle and inherently human concepts, or even hard decision making tasks in challenge situations. We can see intelligent algorithms applied to speech and image recognition, email spam classification, advertising, fraud detection, autonomous self-driving cars, virtual assistants, security, finance and several other applications in our lives. Never in the history we have witness such a great bet in the future of machines.

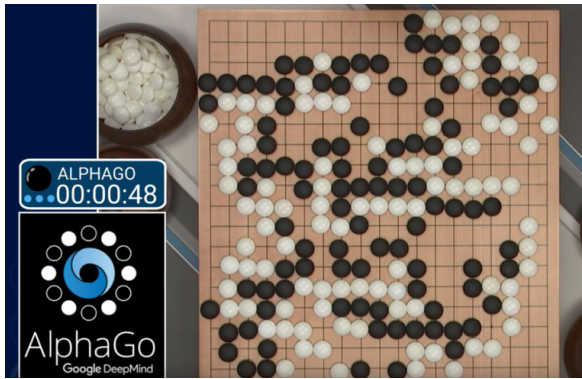
Besides, most recently, the cutting edge achievements in AI have shown algorithms capable of overcome top human intelligence in difficult problems. In 2015, Google DeepMind created an AI that learned how to play 49 Atari games using the same learning algorithm (including the same hyper parameters), using as the input only the pixels from the screen. The algorithm, called Deep Q-Learning (Mnih *et al.*, 2015), was a breakthrough achievement in the mission of accomplishing a general purpose machine learning agent in a wide variety of games.

In the same year, DeepMind left another big mark in the human history. A computer program called AlphaGo defeated the best human Go player for the first time. Go consists of a very complex board game, with more than 10^{170} configurations, and represents one of the biggest challenges to human intelligence and an unconceivable problem to a computer until then. Going even further, in 2017, Silver *et al.* (2017) introduced AlphaGo Zero, an evolution of the previous version, capable of leaning how to play without any outside data from human games, but just playing with itself. Figure 1.1 illustrates these successful examples.

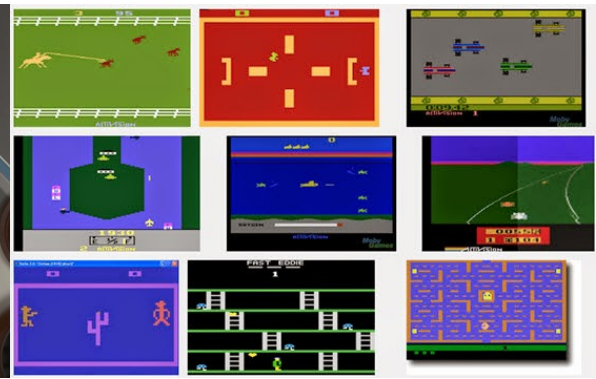
Another big field in Computer Science is mobile robotics, which plays a major role in the future of the industry and the forefront of academic research recently. Robotics can address state of the art challenges in different domains, such as Electronics Engineering, Mechanical Engineering, Control Theory and, of course, Machine Learning.



(a) DQN results in Atari games



(b) AlphaGo against Lee Sedol



(c) Different Atari games

FIGURE 1.1 – DeepMind recent achievements.

AI finds in Robotics several applications, as computer vision, path planning and even locomotion, and in this last one we can find one of its biggest challenges. In Atari and board games we can easily define a goal, win, but how can we define the agility and flexibility of a walk or a jump movement? In this sense, several works have been conducted in the problem of trying to reproduce human movements in humanoid robotics agents, and more specifically, in a simulated scenario.

One of the biggest initiatives of the research community to foster the study of robotics is RoboCup. RoboCup established itself as one of the main international robotics competition in the world, and a powerful scientific conference. It pushes state-of-the-art research in robotics by maintaining different technical challenges with one thing in common, making robots play soccer, with the mission of developing a humanoid robot team capable of win the champion of FIFA World Cup in 2050.




FIGURE 1.2 – Simulated humanoid agent movements and AI.

One RoboCup's particular league is the RoboCup 3D Soccer Simulation League, consisting of a soccer match between two teams, each one composed by up to 11 simulated NAO robots from Aldebaran Robotics. This league address both high level and low level robotics challenges, like path planning and locomotion, and greatly helped improving our understand of human movements like walk and kick the ball.



FIGURE 1.3 – Robocup symbol and Soccer 3D Simulation league match.

1.2 Problem Statement

Inspired by the RoboCup Soccer Simulation 3D (Soccer3D or SS3D) league, this dissertation's objective is to learn a high level soccer behavior for simulated humanoid robots, more specifically, the behavior of kicking the ball towards a planned final distance from the agent. In this sense, the algorithm should input the current game state, including agent's and ball's positions, and output a sequential movement for each robot's joint.

1.3 Approach

Instead of employing a deterministic and off-line built single movement, called Keyframe movement, we intend to use a model-free deep reinforcement learning based approach that tries to learn the desired behavior by interacting with an environment (the simulation server) and receiving rewards depending on which actions it chooses.

1.4 Literature Review

Reinforcement Learning techniques have been in increasingly study in the past 30 years. In this time, some specific works have established famous breakthroughs in this field. For instance, Temporal-Difference (TD) learning algorithms (BARTO *et al.*, 1983) created the groundwork for many future RL algorithms, making use of value function estimation and policy and value iteration methods. In the same way, Q-Learning (WATKINS, 1989) introduced off-policy model-free learning with *Sarsa*, estimating action-value functions and making room to future DRL algorithms.

After these works, Williams (1992) created a new paradigm presenting the REINFORCE algorithm, which introduces policy search methods. This technique estimate the optimal-policy function π_* directly, without the need of value or action-value functions and works way better in continuous action space, as we see in the robotics world. The most recent approaches developed make use and improve this method, such as the Deep Deterministic Policy Gradients (DDPG) algorithm, introduced by (LILLICRAP *et al.*, 2015), the Trust Region Policy Optimization (TRPO) algorithm, presented in (SCHULMAN *et al.*, 2015), and the Proximal Policy Optimization (PPO) algorithm, given in (SCHULMAN *et al.*, 2017).

Another recent breakthrough in the RL field was given by Deep Neural Networks, making possible to escalate the classical algorithms to high dimensional problems. This process was marked by the incredible results of the Deep Q-Networks (DQN) algorithm (MNIH *et al.*, 2015). This technique was able to learn 49 Atari games directly from raw pixels of the screen. It introduces the idea of modeling the action-value function as a neural network, and handle the inherently instability problem from value function approximation by introducing also two techniques: Experience Replay (LIN, 1992) and Target Networks Mnih *et al.* (2015).

Some recent works started to address humanoid movements in the RL field. However, the main problem which comes with that endeavor is that simple rewards functions or naively selected ones, such as the score for the games problems, can lead to results that do not match the expectations. This is the common case in continuous control tasks, like

locomotion. In this sense, (HEESS *et al.*, 2017) proposed that rich and robust behaviors can emerge from simple reward functions, if the environment itself contains sufficient richness and diversity. This work introduces scenarios with a lot of obstacles and varying levels of difficulty, which are presented to the agent as an implicit curriculum, making possible to overcome increasingly hard challenges. A similar idea was also introduced by (BENGIO *et al.*, 2009)

Others DeepMind’s papers introduced methods to learn to imitate human movements, like (MEREL *et al.*, 2017) and (WANG *et al.*, 2017), which combines supervised learning and Generative Adversarial Imitation Learning (GAIL) ((HO; ERMON, 2016)), in a way that accentuates their individual strengths and address their limitations.

Another recent state-of-the-art work in imitation learning is (PENG *et al.*, 2018). This work makes possible for a motion capture actor to supply a set of reference motions for style, and then generate goal-directed and physically realistic behaviors from them. The approach used for this is designing a reward function which combines rewarding motions that resemble reference animation data, and also achieving additional task objectives.

1.5 Contributions

This work’s major contribution is successfully applying deep reinforcement learning (DRL) algorithms in the task of learning a behavior to kick a ball towards a planned final distance from the agent in the Soccer3D environment domain. To the best of our knowledge, this work is the first that makes use of a DRL approach to address this specific challenge.

1.6 Outline of this Dissertation

This dissertation is organized as follows:

- **Chapter 1** introduces this dissertation by describing the motivation behind the problem we address, by the Literature review and by summarizing the contributions.
- **Chapter 2** describes a brief theoretical background of reinforcement learning and neural networks.
- **Chapter ??** describes some experiments using OpenAI gym frameworks

2 Reinforcement Learning

Reinforcement Learning (RL) is an area of machine learning that studies how an agent interacts with an environment receiving instant rewards (SUTTON; BARTO, 1998). The standard RL setup consists of an agent interacting with an environment E in discrete timesteps. At each timestep t , an agent receives an observation S_t , takes action A_t and ends up in state S_{t+1} , receiving a scalar reward R_{t+1} . Figure 2.1 illustrates this notion.

RL is a simple yet powerful framework that defines the interaction between the agent and the environment. The agent learns from the environment in order to achieve long-term goals.

When modelling a problem with RL, there are 3 main components that must be carefully modeled:

- **State space:** set of values that define the state of the agent in the environment. This represents the information available to the agent.
- **Action space:** set of values that define the possible actions the agent can take in the environment.
- **Reward signal:** defines how the agent is rewarded in regards to the goal.

The agent's goal is to maximize the cumulative reward it gets when interacting with the environment. Note that normally these interactions are episodic, which means that




FIGURE 2.1 – Agent interacting with environment.

the task restarts after some number of steps.

The following sections will describe the theory behind RL.

2.1 Markov Decision Process

A Markov Decision Processes (MDP) defines a mathematical framework that formally describes this interaction with the environment. It is said that if the state of the agent contains all the necessary information from the past, it has the Markov Property and can be defined formally as:

$$\mathcal{P}(S_{t+1}|S_t) = \mathcal{P}(S_{t+1}|S_1, S_2, \dots, S_t) \quad (2.1)$$

Therefore, a RL problem is an MDP if it has the Markov Property.

A particular MDP is defined by

- State set, \mathcal{S} .
- Action set, \mathcal{A} .
- Initial state distribution $P(S_1)$.
- One-step transition dynamics of the environment E , denoted as $\mathcal{P}(S_{t+1}|S_t, A_t)$.

In general, the environment may be partially observed so that the entire history of state-action pairs would be required to fully describe the state.

2.2 Return

The return from state S is defined as the sum of discounted future reward from timestep t :

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad (2.2)$$

Notice that $\gamma \in [0, 1]$ introduces the concept of discounting so that the agent prefers to gain instant rewards instead of long term rewards. Also, the return depends on the actions being chosen (therefore on the policy π as well) and may be stochastic.

2.3 Policies

A policy π maps the states to a probability distribution over the actions $\pi : \mathcal{S} \rightarrow \mathcal{P}(\mathcal{A})$. It is represented by Equation (2.3).

$$\pi(a|s) = \mathcal{P}[A_t = a|S_t = s] \quad (2.3)$$

The policy fully describes the behavior of the agent and can also be deterministic formalized as $\mu_\theta : \mathcal{S} \rightarrow \mathcal{A}$ with parameter vector $\theta \in \mathbb{R}^2$ and $A_t = \mu_\theta(S_t)$.

2.4 Value functions

The value functions are estimates of the agent's current return.

- State-value function

$$v_\pi(s) = \mathbb{E}[G_t|S_t = s] \quad (2.4)$$

- Action-value function

$$q_\pi(s, a) = \mathbb{E}[G_t|S_t = s, A_t = a] \quad (2.5)$$

The value function estimates how good the agent's current state is (or how good it is for it to take a given action at that state).

2.5 Goal of RL

The goal of RL is to learn the optimal behavior (described by the policy) of the agent. This policy maximizes the expected return $\mathbb{E}_{R_i, S_i \sim E, A_i \sim \pi}[G_1]$. We will now define optimality for the value function and policy:

2.5.1 Optimal Value Function

The value functions are optimal if they are the highest value for all states (or state-action pairs).

- Optimal State-value function: $v_*(s) = \max_{\pi} v_\pi(s)$.
- Optimal Action-value function: $q_*(s, a) = \max_{\pi} q_\pi(s, a)$.

Therefore, solving an MDP means finding the optimal value function (or optimal policy).

2.5.2 Optimal Policy

Let the partial ordering of policies be

$$\pi \geq \pi' \text{ if } v_\pi(s) \geq v_{\pi'}(s)$$

The optimal policy is better than any other, $\pi_* \geq \pi, \forall \pi$ and achieves the optimal value functions.

An interesting property of the optimal policy is that it can be implicitly defined with q_* . The optimal policy can be defined as the following deterministic policy

$$\pi_*(a|s) = \begin{cases} 1, & \text{if } a = \underset{a'}{\operatorname{argmax}} q_*(s, a') \\ 0, & \text{otherwise} \end{cases}$$


2.6 Function Approximation

Many classic tabular RL algorithms work by estimating the value functions for each state, for example, Value Iteration methods. A problem arises when the state/action spaces are extremely large, or if the state/action spaces are continuous. Some examples of large discrete state spaces are board games. This is known as the *curse of dimensionality* (BENGIO *et al.*, 2012). Backgammon has over 10^{20} board states (LEVNER, 1976) and the game of Go (with board size of 19x19) has over 10^{170} states (TROMP; FARNEBÅCK, 2016). Figure 2.2 illustrates these games.

There are 2 problems of having so many states in realistic situations:

- It will be impossible to store all the states in memory.
- It will be impossible to visit all the states in training.

For continuous domains, we have an analogous situation since it is possible to discretize the state or action space and then apply tabular methods. Figure 2.3 illustrates 3 continuous control problems and the number of states resulting from one possible discretization.





(a) Game of Backgammon.



(b) Game of Go.

FIGURE 2.2 – Example of board games that have large discrete state spaces.

Crawler
 10^2 Hopper
 10^{10} Humanoid
 10^{100} FIGURE 2.3 – Sample control tasks from the (BROCKMAN *et al.*, 2016) toolkit.

If the discretized space is large, we still have the problem described above. An approach to deal with this problem is to use a function approximation that estimates the value function. The approximate value function, $\hat{v}(s, \theta)$ will be defined by a parameter weight vector $\theta \in \mathbb{R}^n$ and

$$\hat{v}(s, \theta) \approx v_{\pi}(s)$$

The interesting part of function approximation is that this function estimator can be any function, ranging from a simple linear function of the states to a complex NNs. Normally the number of weights is much less than the number of states.

In this work, we focused on using neural networks as function approximators; Chapter 3 lays the groundwork for the theory regarding Neural Networks.

2.7 RL algorithms

There are 3 classes of RL algorithms: methods based on value functions, methods based on policy search and a hybrid approach which combines value functions and policy search (SUTTON; BARTO, 1998).

2.7.1 Value Function Methods

Value function methods are based on estimating v or q functions. Monte-Carlo (MC) methods and TD Learning are the two methods generally used for predicting the value function (SUTTON; BARTO, 1998).

2.7.1.1 Monte-Carlo

MC prediction are model-free (i.e. they do not require knowledge of the underlying dynamics) methods that learn directly from complete episodes of experience. They estimate the value functions by averaging sample returns and therefore can only be applied to episodic MDPs.

2.7.1.2 Temporal-Difference Learning

TD Learning are model-free and can learn from incomplete episodes. TD Learning updates the estimate by *bootstrapping*, which means it updates the value function using an existing estimate.

The simplest TD method, known as TD(0), updates the value function as

$$V(S_t) = V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)] \quad (2.6)$$

TD has some advantages in comparison to MC: TD can learn before the end of that episode (or in non episodic MDPS) and have low variance and some bias. MC on the other hand has high variance, zero bias, having good convergence properties (even with function approximation).

2.7.2 Policy Search Methods

Policy search methods do not have an estimate of the value function. These algorithms typically maintain a parameterized policy π_θ whose parameters are updated to maximize

the expected return. The Policy Gradient Theorem defines that for any differentiable policy $\pi_\theta(s, a)$ the policy gradient is

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta}[\nabla_{\pi_\theta} \log \pi_\theta(s, a) Q^{\pi_\theta}(s, a)] \quad (2.7)$$

An example usage of this theorem is seen in the REINFORCE algorithm that uses Equation (2.7) to update the policy's weight using gradient ascent (WILIAMS, 1992).

2.7.3 Actor-Critic Methods

Actor-critic methods also work by updating a parameterized policy π_θ but instead use a *critic* to estimate the action-value function Q^{π_θ} ,

$$Q_w(s, a) \approx Q^{\pi_\theta}(s, a)$$

where w is the weight vector that parameterizes Q . Therefore actor-critic algorithms maintain two sets of parameters: the *critic* that updates that action-value function parameters w and the *actor* that updates the policy parameters θ in the direction suggested by the *critic*.

Figure 2.4 illustrates the architecture of actor-critic algorithms.




FIGURE 2.4 – Actor-critic architecture from (SUTTON; BARTO, 1998).

2.8 Curriculum Learning

Reinforcement learning has been used to solve non-trivial tasks in locomotion (SCHULMAN *et al.*, 2015) and video games (MNIH *et al.*, 2015).

However, there are many tasks that are hard to design a reward function that is easily maximized and that when optimized yields the expected behavior (FLORENSA *et al.*, 2017). For example, tasks with sparse rewards are still challenging since it is very hard to achieve the goal following random exploration (MATHISEN *et al.*, 2017). One approach to overcome this problem is to use curriculum learning (BENGIO *et al.*, 2009).

Curriculum learning is a sequence of training criteria where tasks are ordered by increasing difficulty and training moves from to a harder task only when it has mastered the easier tasks. Usually, to apply curriculum learning to a task, the researcher must be able to order the subtasks by difficulty and decide the heuristics for when a subtask has been mastered (MATHISEN *et al.*, 2017).

Curriculum learning is a recent idea that has been applied to a wide range of RL problems. An interesting result is Zaremba e Sutskever (2014) that applies a variant of curriculum learning to evaluate short computer programs. Some other examples include (BANSAL *et al.*, 2017; MATHISEN *et al.*, 2017).

3 Deep Learning

In the last chapter we described the mathematical background behind reinforcement learning and how a problem can be modeled with it. We also saw that for problems with continuous state and/or action spaces, it can be necessary to use function approximators such as neural networks. In this chapter we will describe the theoretical background of NNs, as function approximators.

Deep learning is a field of Machine Learning that in part studies deep neural networks and how they learn.

3.1 History

Today, Deep learning is a very significant topic for the scientific community yet its history goes back since the 1940s (GOODFELLOW; COURVILLE, 2016).

There has been a major resurgence of deep learning mainly because computers today became faster. Training large deep learning models are computationally very expensive and could be sped up with the use of graphics processing units (GPUs).

Artificial Neural Networks are one of the earliest learning algorithms. The novel concept behind it was to create mathematical models that tried to mimic the brain. NNs were based on the human brain's neurological structure, more specifically neurons and how they connect with other neurons. Figure 3.1 illustrates a single neuron cell.

A neuron cell is capable of transmitting information by sending electrical and chemical signals through axons and they connect to other neurons forming neural networks. The human brain, for example, has billions of connected neurons. NNs, therefore, are a simplified mathematical representation of neurons: essentially, a neuron can be viewed as a cell that receives one or more inputs and produces one or more outputs.

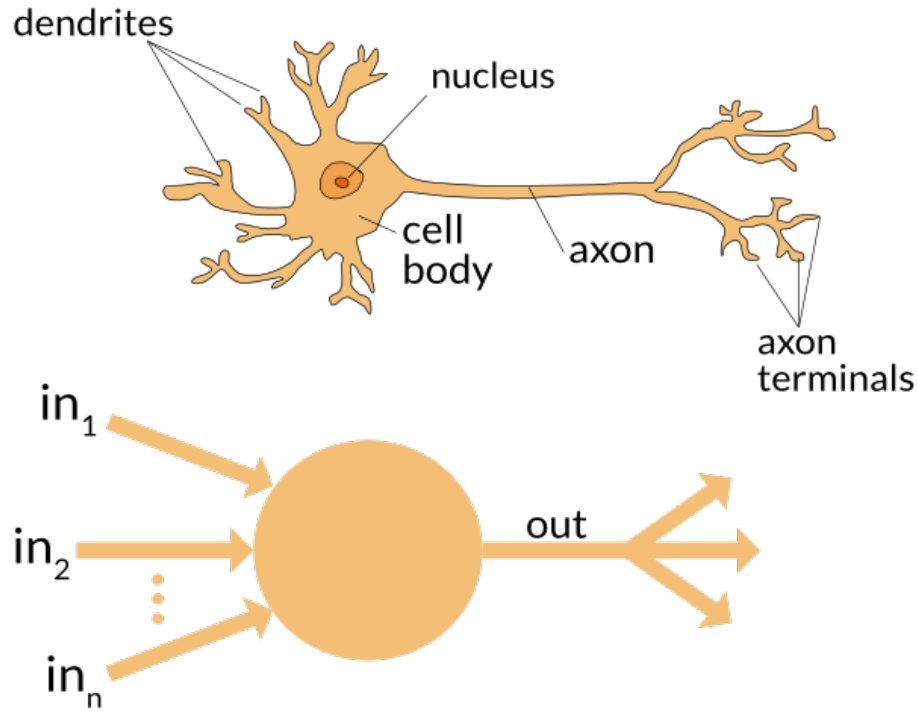


FIGURE 3.1 – Neuron unit cell structure.

3.2 Neural Networks

3.2.1 Representation

In simple terms, a neural network is a mathematical structure that receives an input and calculates an output. We will describe a simple NN known as feedforward neural network or Multi Layer Perceptron (MLP).

A neural network is composed of multiple layers where each layer has one or more nodes (neurons). The NN receives an input and computes the output according to the network architecture and parameters θ . Each input, called sample or example, is a n -dimensional feature vector x , and is represented as a column vector:

$$x = \begin{bmatrix} x_1 & x_2 & x_3 & x_4 & \dots & x_{n-1} & x_n \end{bmatrix}^T \quad (3.1)$$

x is known as the input layer, and each coordinate represents each of the sample's features. The output is a k -dimensional vector y , known as the output layer. Every other layer that is not the input nor the output layer is known as a hidden layer.

In this work, note that every time we refer to an element $z^{[l]}$ we are referring to an element of the l -th layer.

Figure 3.2 illustrates a very simple NN with a 3-dimensional feature vector, 1-dimensional

output vector and only one hidden layer.

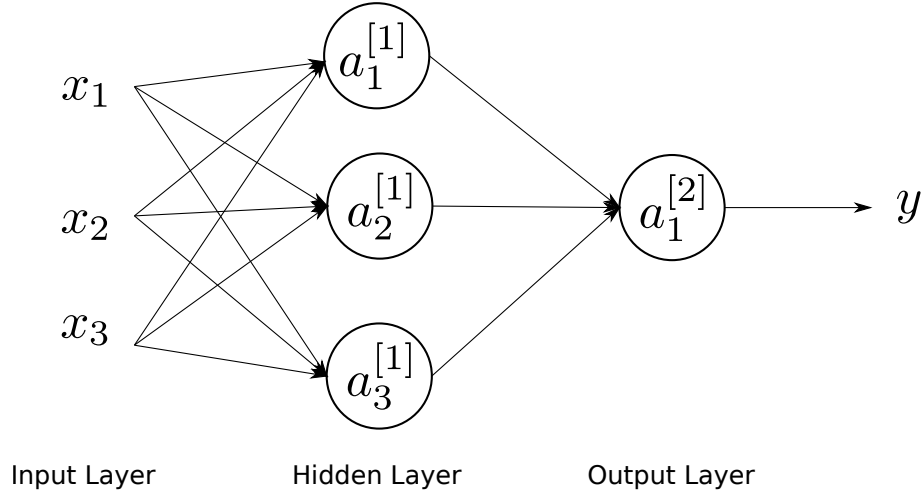


FIGURE 3.2 – Simple neural network.

Every node of the network is computed using the values from the previous layer. For the MLP, the node computation model is done through a linear model $f(x; \theta)$, with θ consisting of internal parameters w (weights) and b (bias). w has the same dimensions as the input and b is a scalar. In this work, we treat w as a column vector.

$$w = \begin{bmatrix} w_1 & w_2 & w_3 & w_4 & \dots & w_{n-1} & w_n \end{bmatrix}^T \quad (3.2)$$

The model is defined as

$$z = w^T x + b \quad (3.3)$$

The output of the node a , also known as the activation, is calculated with x , b and a function $\sigma : \mathbb{R} \rightarrow \mathbb{R}$:

$$a = \sigma(z) = \sigma(w^T x + b) = \sigma\left(\sum_{i=1}^n w_i x_i + b\right) \quad (3.4)$$

Figure 3.3 illustrates the output of a single neuron.

σ is called activation function, applied to the model output (z). The activation function normally is a non-linear function and its most used common choice today is the **rectified linear unit** (NAIR; HINTON, 2010) or ReLU and is defined as $\sigma(z) = \max\{z, 0\}$, shown in 3.4.

Other examples of popular activation functions are $\sigma(z) = \frac{1}{1+e^{-z}}$, known as the sigmoid function and $\tanh(z)$ (XU; LI, 2016).




FIGURE 3.3 – Node computation.




FIGURE 3.4 – Linear rectified unit plot.

A neural network is therefore a connection of multiple neurons as inputs to other neurons. Let us return our attention to the network from Figure 3.2. We denote $(W, b) = (W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]})$ as the network parameters, where $W^{[l]}$ is the matrix formed by concatenating the individual parameters of the l -th layer such that W_{ij} denotes the parameter associated with the connection between neuron j from the l -th layer and neuron i from layer $l + 1$. Similarly, $b^{[l]}$ is the concatenation of the biases b of each neuron of layer $l + 1$ and $b_i^{[l]}$ the bias of the i -th neuron of layer $l + 1$.

For the network depicted in Figure 3.2, we are now capable of calculating the output y .

For the first layer:

$$z^{[1]} = W^{[1]}x + b^{[1]} \quad (3.5)$$

$$a^{[1]} = \sigma(z^{[1]}) \quad (3.6)$$

And for the second layer:

$$z^{[2]} = W^{[2]}a^{[1]} + b^{[2]} \quad (3.7)$$

$$y = a^{[2]} = \sigma(z^{[2]}) \quad (3.8)$$

We can think of the neural network as a recursive structure where each activation is calculated with the past's layers activation. $a^{[l]}$ can be computed as:

$$z^{[l]} = W^{[l]}a^{[l-1]} + b^{[l]} \quad (3.9)$$

$$a^{[l]} = \sigma(z^{[l]}) \quad (3.10)$$

Note also that $a^{[1]} = x$.

3.2.2 Vectorization

In the previous section, we focused on computing the output of the neural network given a single sample x . In modern days, however, we are normally interested in calculating the output of a neural network for thousands or even millions of samples.

For m samples, we could naively compute the output of the neural network for each example $x^{(i)}$ with Algorithm 1.

Algorithm 1: Naive algorithm for computing NN output of m samples.

Result: Output of NN for m samples.

```

for  $i = 1$  to  $m$  do
     $a^{[0]} = x^{(i)}$ 
    for  $j = 1$  to  $L$  do
         $z^{[j](i)} = W^{[j]}a^{[j-1](i)} + b^{[j]}$ 
         $a^{[j](i)} = \sigma(z^{[j](i)})$ 
    end
     $y^{(i)} = a^{[L](i)}$ 
end

```

In practice this algorithm runs relatively slow since it computes the output of the

network sequentially for each sample and we will apply vectorization to achieve a much faster algorithm.

Firstly, Vectorization is a technique that transforms a set of computations done sequentially in a for loop into matrix operations. We will now vectorize our initial problem.

Let us define the following matrices

- X is a matrix where column i is the i -th sample $x^{(i)}$ and $W^{[l]}$. We can analogously define matrices A and Z . $X = \begin{bmatrix} x^{(1)} & x^{(2)} & \dots & x^{(m)} \end{bmatrix}$.
- $W^{[l]}$ is the weight matrix for the l -layer, as defined in Section 3.2.1.
- $b^{[l]}$ is the bias vector for layer l , as defined in Section 3.2.1.

For the vectorized version of Algorithm 2.

Algorithm 2: Vectorized algorithm for computing NN output of m samples.

Result: Output of NN for m samples.

```

for  $j = 1$  to  $L$  do
    |  $Z^{[j]} = W^{[j]}A^{[j-1]} + b^{[j]}$ 
    |  $A^{[j]} = \sigma(Z^{[j]})$ 
end

```

The vectorized version is much better since these matrices operations can be greatly sped up on GPUs or hardware that have support for Streaming SIMD Extensions (SSE). The great majority of popular deep learning libraries use as much vectorization as they can.

3.2.3 Deep Neural Networks

We can extend the architecture we saw on Section 3.2.1 to a higher number of layers called deep neural networks. Figure 3.5 illustrates a neural network with multiple hidden layers.

The activation of layer $l + 1, l > 1$, can be calculated according to Equation 3.9.

3.3 Learning

As described in the previous section, feedforward networks defines a mapping $y = f(x; \theta)$, or, equivalently, serve as general function approximators.




FIGURE 3.5 – Neural network with multiple layer.

This section describes deep learning algorithms that learn the value of the parameters θ that best approximates f .

Most deep learning algorithms need to describe a cost function and optimization algorithm.

3.3.1 Cost Function

The cost function $J(\theta)$ is a metric of how good our estimator is to our dataset: the lower the cost function, the less error the model has when predicting y . It describes the function that we wish to minimize and normally it involves an average of the errors between the the target value of a sample and the predicted value of the estimator for the sample.

For example, for the problem of linear regression where X are the sample values and y are the target values for our dataset, we wish to learn the parameters θ from our function approximator f . The most used cost function for this problem is

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m (y_i - f(x_i, \theta))^2 \quad (3.11)$$

Another example is the logistic regression problem. The most common cost function

for this problem is

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y_i \log f(x_i, \theta) + (1 - y_i) \log 1 - f(x_i, \theta)] \quad (3.12)$$

Choosing a good cost function greatly impacts how good the design of the deep neural network is.

3.3.2 Optimization Algorithms

After defining a cost function, finding the parameters that effectively minimize the cost function is not trivial. The non linearity of NN causes most cost functions to be non-convex and therefore, neural networks are usually trained by iterative, gradient-based optimizers. In practice, however, this is very unlikely (SWIRSZCZ *et al.*, 2017; GOODFELLOW *et al.*, 2014; LIN *et al.*, 2017).

Section 3.4 describes how neural networks gradients can be calculated.

3.3.2.1 Gradient Descent

One of the most common optimization algorithm is Gradient Descent. It is an iterative first-order method that tries to find a local minimum for a function f by moving in the negative direction of its gradient. Specifically for neural networks, our goal is to minimize the cost function $J(\theta)$. The algorithm's update step is Equation (3.13) done for each θ_i

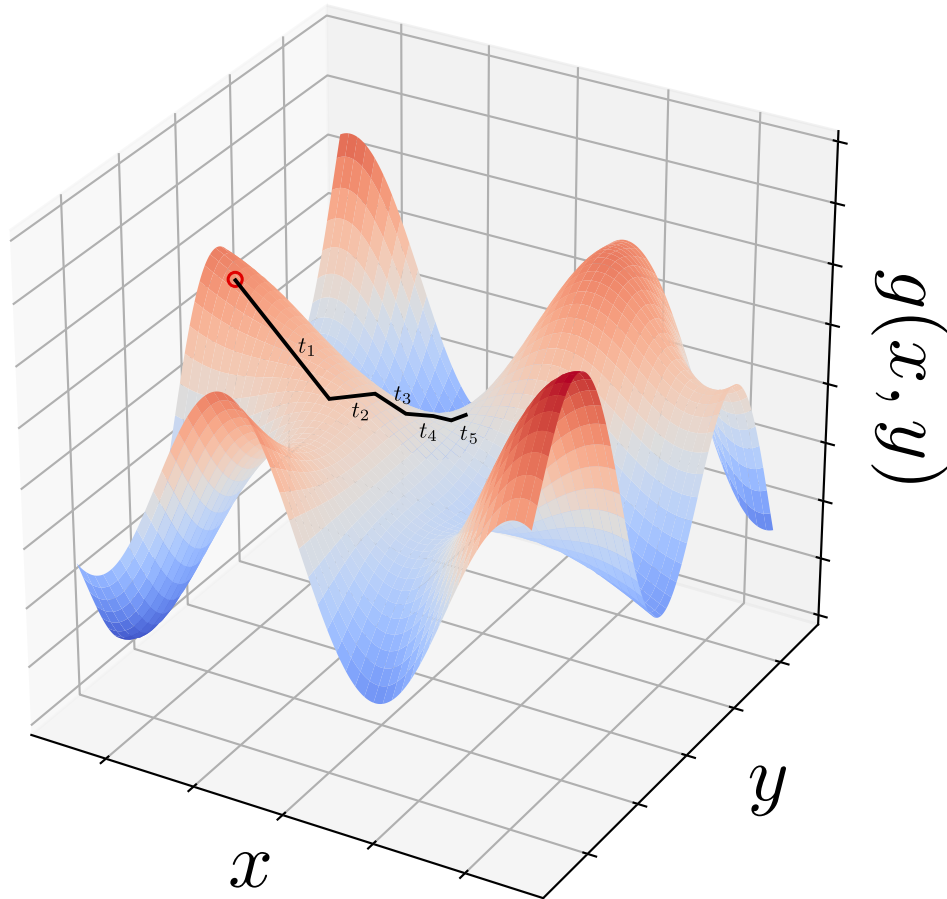
$$\theta_i = \theta_i - \alpha \frac{\partial J}{\partial \theta_i} \quad (3.13)$$

The learning rate, α measures how big the update step in the direction of the gradient will be.

For example, let $g : \mathbb{R}^2 \rightarrow \mathbb{R}$ given by $g(x, y) = y \sin(x) - x \cos(y)$. Figure 3.6 illustrates 5 steps of the gradient descent algorithm. Notice how the red circle is the starting point and after the last iteration, the algorithm arrives approximately at a local minimum.

3.3.2.2 Stochastic Gradient Descent

Another popular algorithm is Stochastic Gradient Descent (SGD) that is a stochastic approximation of Gradient Descent. Instead of calculating $J(\theta)$ as an average between all the samples, the gradient update step is done once per sample. It calculates an approximation of the gradient and updates the parameters θ_i with every sample. Algorithm 3

FIGURE 3.6 – Gradient descent on function g .

describes SGD

Algorithm 3: SGD algorithm.

Result: Function minimum

while *An approximate minimum is not obtained* **do**

 Randomly shuffle examples in the training set

for $i = 1, 2, \dots, m$ **do**

$\theta_i = \theta_i - \alpha \nabla J_i(\theta)$

end

end

The gradient approximation for SGD is much faster to compute, however, since it is a noisy estimate of the gradient, SGD may need more steps to converge to a local minimum.

3.3.2.3 Mini-Batch Gradient Descent

Finally, Mini-Batch Gradient Descent combines the idea from Gradient Descent and SGD. It computes the gradient using more than one training example, called a *mini-batch*, at each step. It may result in smoother convergence and the code can be accelerated by

making use of vectorization.

3.4 Back Propagation

We have seen how a few gradient-based optimization algorithms work. To minimize the cost function we need to compute the partial derivatives of the cost function w.r.t. to each weight of the neural network. Calculating these partial derivatives manually would be computationally infeasible for large networks. In deep neural networks, by using the chain rule, we can calculate the partial derivatives in an efficient and compact way.

Backpropagation is an algorithm that computes the gradients of the loss function w.r.t. each input from the network.

Bibliography

BANSAL, T.; PACHOCKI, J.; SIDOR, S.; SUTSKEVER, I.; MORDATCH, I. Emergent complexity via multi-agent competition. **CoRR**, abs/1710.03748, 2017. Disponível em: <<http://arxiv.org/abs/1710.03748>>.

BARTO, A. G.; SUTTON, R. S.; ANDERSON, C. W. Neuronlike adaptive elements that can solve difficult learning control problems. 1983.

BENGIO, Y.; COURVILLE, A. C.; VINCENT, P. Unsupervised feature learning and deep learning: A review and new perspectives. **CoRR**, abs/1206.5538, 2012. Disponível em: <<http://arxiv.org/abs/1206.5538>>.

BENGIO, Y.; LOURADO, J.; COLLOBERT, R.; WESTON, J. Curriculum learning. In: **Proceedings of the 26th Annual International Conference on Machine Learning**. New York, NY, USA: ACM, 2009. (ICML '09), p. 41–48. ISBN 978-1-60558-516-1. Disponível em: <<http://doi.acm.org/10.1145/1553374.1553380>>.

BROCKMAN, G.; CHEUNG, V.; PETTERSSON, L.; SCHNEIDER, J.; SCHULMAN, J.; TANG, J.; ZAREMBA, W. **OpenAI Gym**. 2016. Acesso em: 10 maio de 2018.

FLORENSA, C.; HELD, D.; WULFMEIER, M.; ABBEEL, P. Reverse curriculum generation for reinforcement learning. **CoRR**, abs/1707.05300, 2017. Disponível em: <<http://arxiv.org/abs/1707.05300>>.

GOODFELLOW, I. J.; VINYALS, O.; SAXE, A. M. Qualitatively characterizing neural network optimization problems. **arXiv preprint arXiv:1412.6544**, 2014.

GOODFELLOW, Y. B. I.; COURVILLE, A. **Deep Learning**. MIT Press, 2016. Disponível em: <<http://www.deeplearningbook.org>>.

HEESS, N.; TB, D.; SRIRAM, S.; LEMMON, J.; MEREL, J.; WAYNE, G.; TASSA, Y.; EREZ, T.; WANG, Z.; ESLAMI, S. M. A.; RIEDMILLER, M.; SILVER, D. Emergence of locomotion behaviours in rich environments. july 2017.

HO, J.; ERMON, S. Generative adversarial imitation learning. In: **Advances in Neural Information Processing Systems**. [S.l.: s.n.], 2016. p. 4565–4573.

LEVNER, D. **Is Brute Force Backgammon Possible?** 1976. Disponível em: <www.bkgm.com/articles/Levner/BruteForceBackgammon/>.

- LILICRAP, T. P.; HUNT, J. J.; PRITZEL, A.; HEESS, N.; EREZ, T.; TASSA, Y.; SILVER, D.; WIERSTRA, D. Continuous control with deep reinforcement learning. **CoRR**, abs/1509.02971, 2015. Disponível em: <<http://arxiv.org/abs/1509.02971>>.
- LIN, H. W.; TEGMARK, M.; ROLNICK, D. Why does deep and cheap learning work so well? **Journal of Statistical Physics**, Springer, v. 168, n. 6, p. 1223–1247, 2017.
- LIN, L.-J. Self-improving reactive agents based on reinforcement learning, planning and teaching. 1992.
- MATHISEN, T.; OLIVER, A.; COHEN, T.; SCHULMAN, J. Teacher-student curriculum learning. **CoRR**, abs/1707.00183, 2017. Disponível em: <<http://arxiv.org/abs/1707.00183>>.
- MEREL, J.; TASSA, Y.; TB, D.; SRINIVASAN, S.; LEMMON, J.; WANG, Z.; WAYNE, G.; HEESS, N. Learning human behaviors from motion capture by adversarial imitation. july 2017.
- MNIH, V.; KAVUKCUOGLU, K.; SILVER, D.; RUSU, A. A.; VENESS, J.; BELLEMARE, M. G.; GRAVES, A.; RIEDMILLER, M.; FIDJELAND, A. K.; OSTROVSKI, G.; PETERSEN, S.; BEATTIE, C.; SADIK, A.; ANTONOGLOU, I.; KING, H.; KUMARAN, D.; WIERSTRA, D.; LEGG, S.; HASSABIS, D. Human-level control through deep reinforcement learning. **Nature**, Nature Publishing Group, a division of Macmillan Publishers Limited. All Rights Reserved., v. 518, n. 7540, p. 529–533, Feb 2015. ISSN 0028-0836. Letter. Disponível em: <<http://dx.doi.org/10.1038/nature14236>>.
- NAIR, V.; HINTON, G. E. Rectified linear units improve restricted boltzmann machines. In: **Proceedings of the 27th International Conference on International Conference on Machine Learning**. USA: Omnipress, 2010. (ICML'10), p. 807–814. ISBN 978-1-60558-907-7. Disponível em: <<http://dl.acm.org/citation.cfm?id=3104322.3104425>>.
- PENG, X. B.; ABBEEL, P.; LEVINE, S.; PANNE, M. V. D. Deepmimic: Example-guided deep reinforcement learning of physics-based character skills. April 2018.
- SCHULMAN, J.; LEVINE, S.; MORITZ, P.; JORDAN, M. I.; ABBEEL, P. Trust region policy optimization. **CoRR**, abs/1502.05477, 2015. Disponível em: <<http://arxiv.org/abs/1502.05477>>.
- SCHULMAN, J.; WOLSKI, F.; DHARIWAL, P.; RADFORD, A.; KLIMOV, O. Proximal policy optimization algorithms. **CoRR**, abs/1707.06347, 2017. Disponível em: <<http://arxiv.org/abs/1707.06347>>.
- SILVER, D.; SCHRITTWIESER, J.; SIMONYAN, K.; ANTONOGLOU, I.; HUANG, A.; GUEZ, A.; HUBERT, T.; BAKER, L.; LAI, M.; BOLTON, A.; CHEN, Y.; LILICRAP, T.; HUI, F.; SIFRE, L.; DRIESSCHE, G. van den; GRAEPEL, T.; HASSABIS, D. Mastering the game of go without human knowledge. **Nature**, Macmillan Publishers Limited, part of Springer Nature. All rights reserved., v. 550, out. 2017.

SUTTON, R. S.; BARTO, A. G. **Introduction to Reinforcement Learning**. 1st. ed. Cambridge, MA, USA: MIT Press, 1998. ISBN 0262193981.

SWIRSZCZ, G.; CZARNECKI, W. M.; PASCANU, R. Local minima in training of neural networks. **stat**, v. 1050, p. 17, 2017.

TROMP, J.; FARNEBÄCK, G. Combinatorics of go. 2016. Disponível em: <<https://tromp.github.io/go/gostate.pdf>>.

WANG, Z.; MEREL, J.; REED, S.; WAYNE, G.; FREITAS, N. de; HEESS, N. Robust imitation of diverse behaviors. july 2017.

WATKINS, C. J. C. H. **Learning from Delayed Rewards**. Tese (Doutorado) — King's College, 1989.

WILIAMS, R. J. Simple statistical gradient-following algorithms for connectionist reinforcement learning. 1992.

XU, R. H. B.; LI, M. Revise saturated activation functions. 2016.

ZAREMBA, W.; SUTSKEVER, I. Learning to execute. **CoRR**, abs/1410.4615, 2014. Disponível em: <<http://arxiv.org/abs/1410.4615>>.

Appendix A - Tipos de Dilema Linear

A.1 Uma Primeira Seção para o Apêndice

A matriz de Dilema Linear M e o vetor de torques inerciais b , utilizados na simulação são calculados segundo a formulação abaixo:

$$M = \begin{bmatrix} M_{11} & M_{12} & M_{13} \\ M_{21} & M_{22} & M_{23} \\ M_{31} & M_{32} & M_{33} \end{bmatrix} \quad (\text{A.1})$$



FIGURE A.1 – Uma figura que está no apêndice

Annex A - Exemplo de um Primeiro Anexo

A.1 Uma Seção do Primeiro Anexo

Algum texto na primeira seção do primeiro anexo.

FOLHA DE REGISTRO DO DOCUMENTO

1. CLASSIFICAÇÃO/TIPO TC	2. DATA 4 de março de 2015	3. DOCUMENTO Nº DCTA/ITA/DM-018/2015	4. Nº DE PÁGINAS 42
5. TÍTULO E SUBTÍTULO: Simulated Humanoid Robot Control with Reinforcement Learning			
6. AUTOR(ES): Luis Guilherme Gomes Aguiar			
7. INSTITUIÇÃO(ÕES)/ÓRGÃO(S) INTERNO(S)/DIVISÃO(ÕES): Instituto Tecnológico de Aeronáutica – ITA			
8. PALAVRAS-CHAVE SUGERIDAS PELO AUTOR: Cupim; Cimento; Estruturas			
9. PALAVRAS-CHAVE RESULTANTES DE INDEXAÇÃO: Cupim; Dilema; Construção			
10. APRESENTAÇÃO: (X) Nacional () Internacional ITA, São José dos Campos. Curso de Mestrado. Programa de Pós-Graduação em Engenharia Aeronáutica e Mecânica. Área de Sistemas Aeroespaciais e Mecânica. Orientador: Prof. Dr. Adalberto Santos Dupont. Coorientadora: Profª. Drª. Doralice Serra. Defesa em 05/03/2015. Publicada em 25/03/2015.			
11. RESUMO:			
12. GRAU DE SIGILO: (X) OSTENSIVO () RESERVADO () SECRETO			